

# Introduction to Deep Learning

## 7. Multilayer Perceptron

**MGMT 735**

**Slides From Mu Li and Alex Smola**

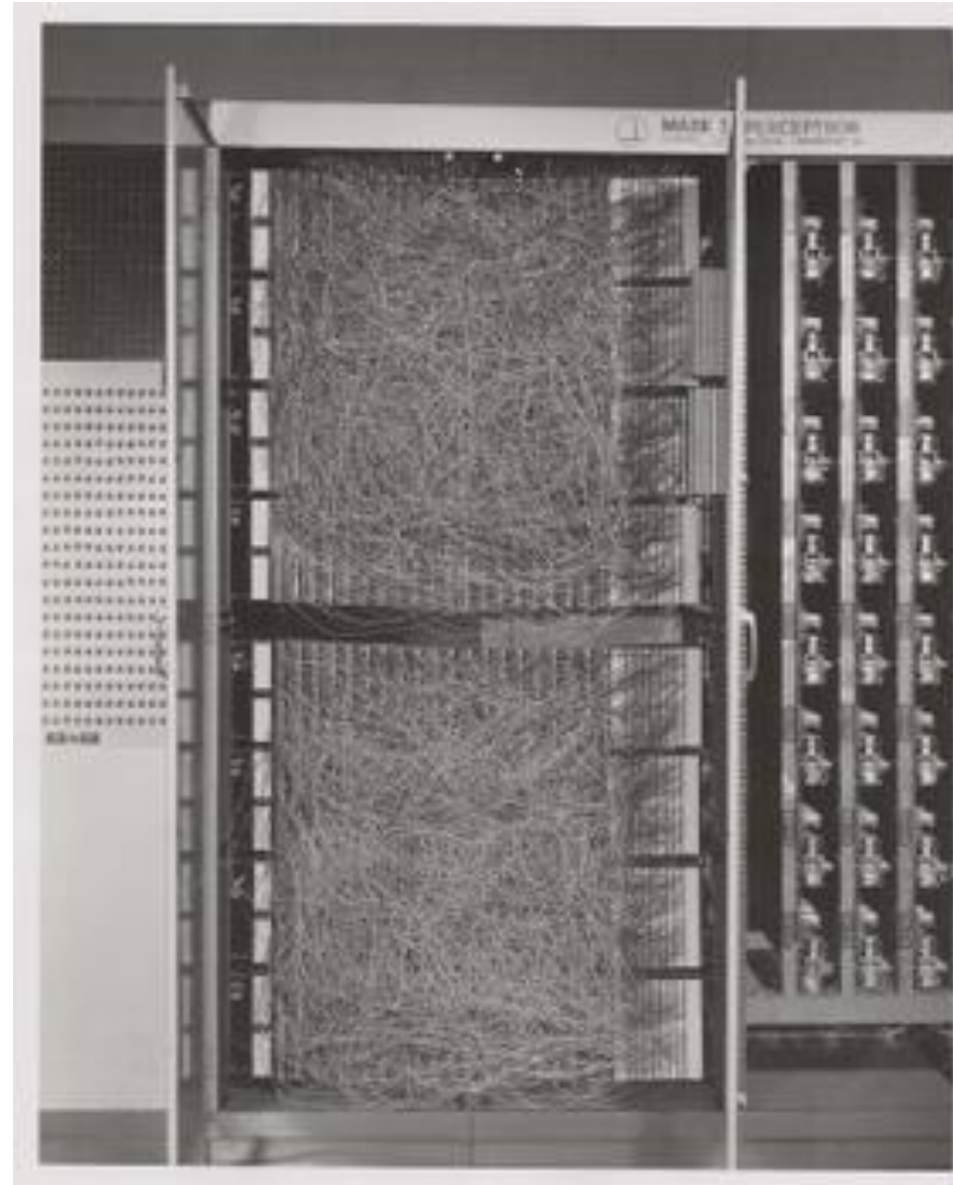
**[courses.d2l.ai/berkeley-stat-157](https://courses.d2l.ai/berkeley-stat-157)**

# Outline

- **Single Layer Perceptron**
  - Decision Boundary
  - XOR is hard
- **Multilayer Perceptron**
  - Layers
  - Nonlinearities
  - Computational Cost

# Perceptron

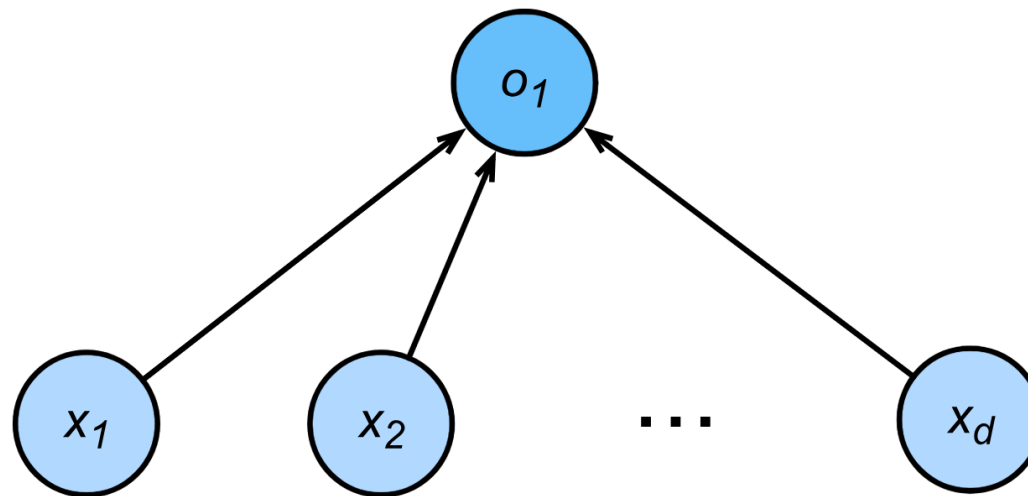
Mark I Perceptron, 1960  
([wikipedia.org](https://en.wikipedia.org/wiki/Mark_I_Perceptron))



# Perceptron

- Given input  $\mathbf{x}$ , weight  $\mathbf{w}$  and bias  $b$ , perceptron outputs:

$$o = \sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b) \quad \sigma(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

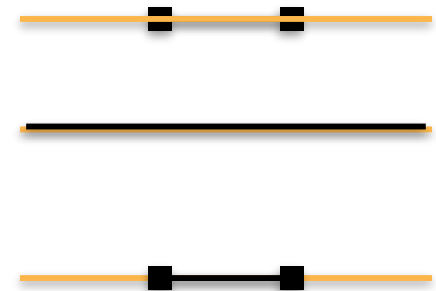


# Perceptron

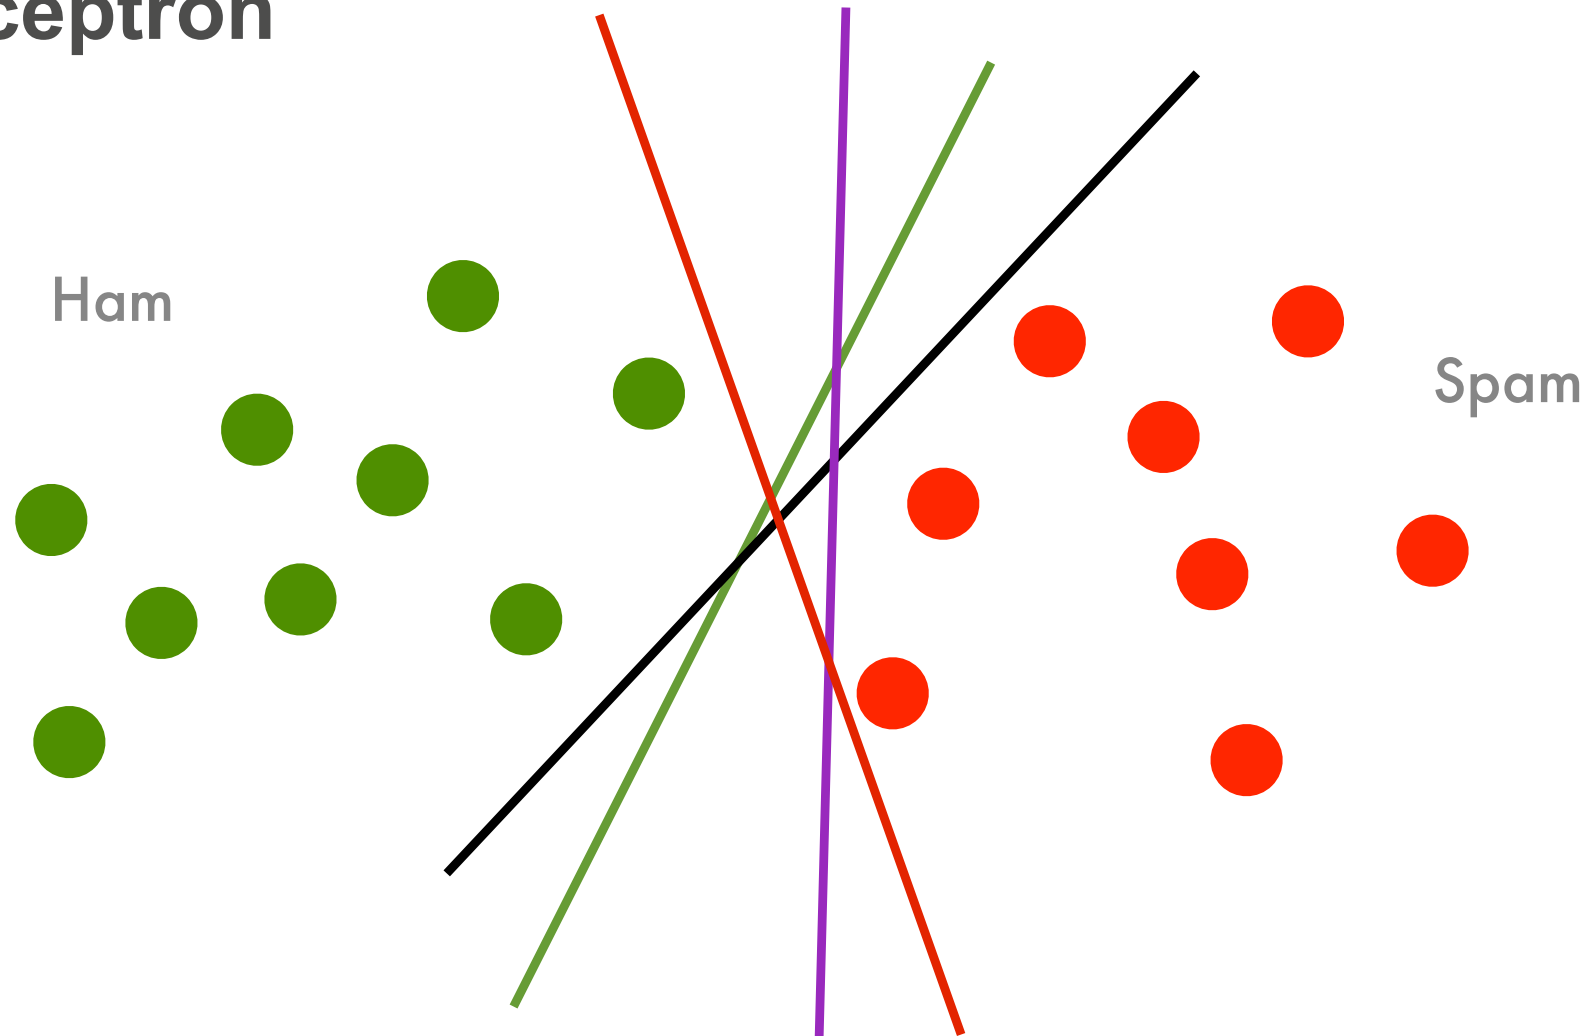
- Given input  $\mathbf{x}$ , weight  $\mathbf{w}$  and bias  $b$ , perceptron outputs:

$$o = \sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b) \quad \sigma(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

- Binary classification (0 or 1)
  - Vs. scalar real value for regression
  - Vs. probabilities for logistic regression



# Perceptron



## Training the Perceptron

**initialize**  $w = 0$  and  $b = 0$

**repeat**

**if**  $y_i [\langle w, x_i \rangle + b] \leq 0$  **then**

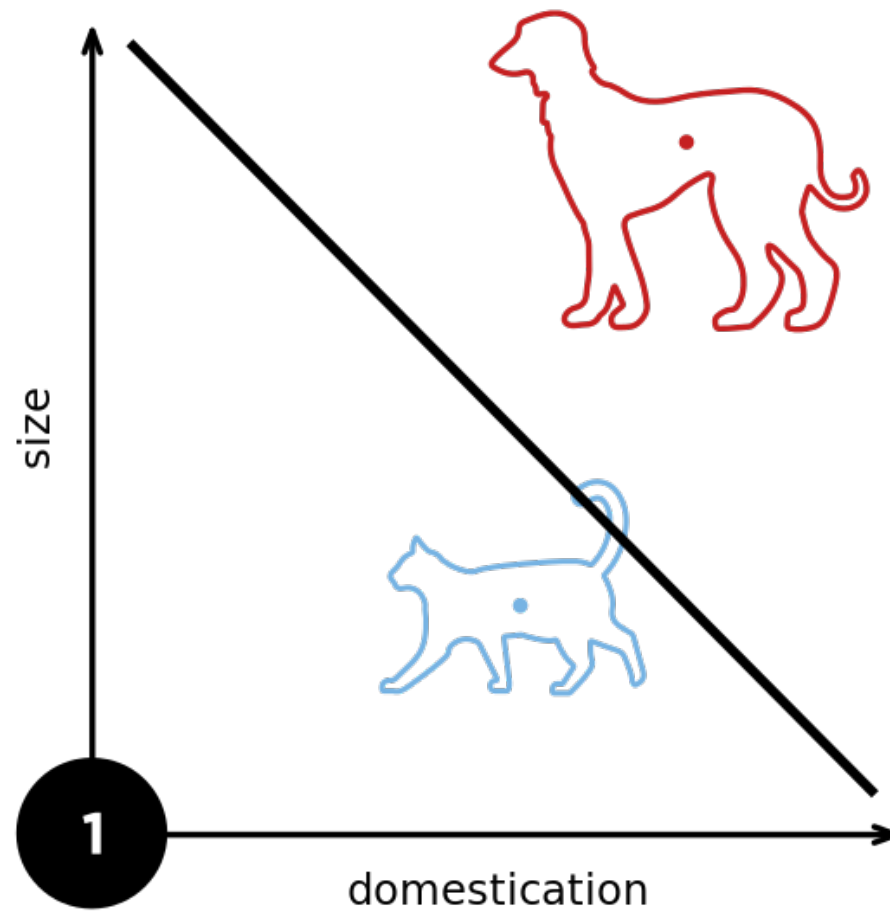
$w \leftarrow w + y_i x_i$  and  $b \leftarrow b + y_i$

**end if**

**until** all classified correctly

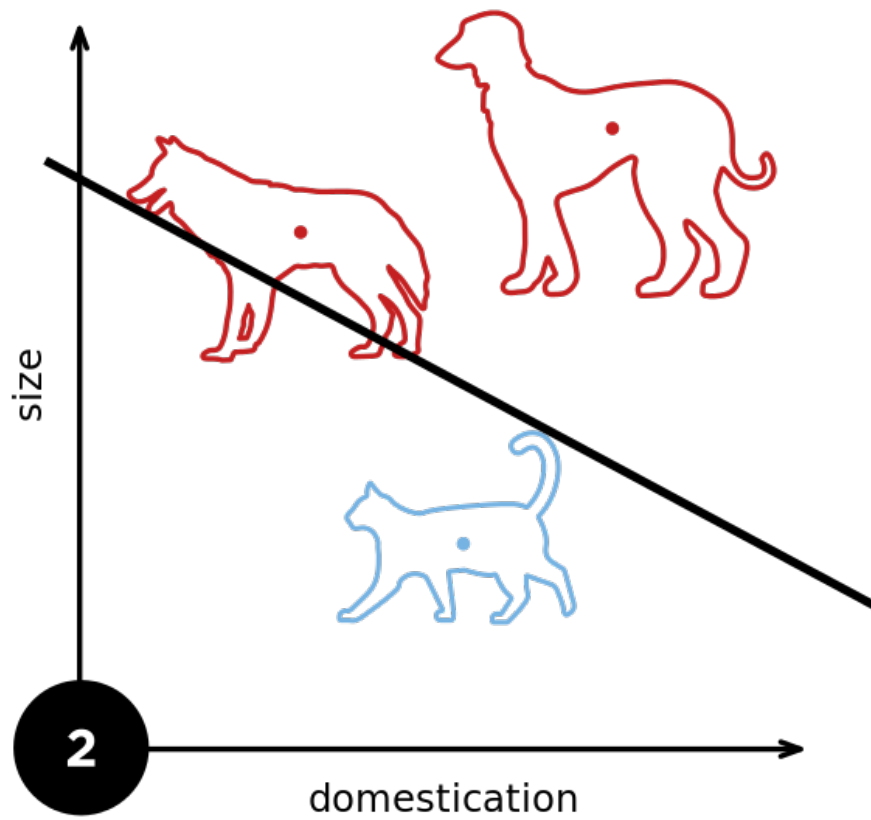
Equals to SGD (batch size is 1) with the following loss

$$\ell(y, \mathbf{x}, \mathbf{w}) = \max(0, -y\langle \mathbf{w}, \mathbf{x} \rangle)$$

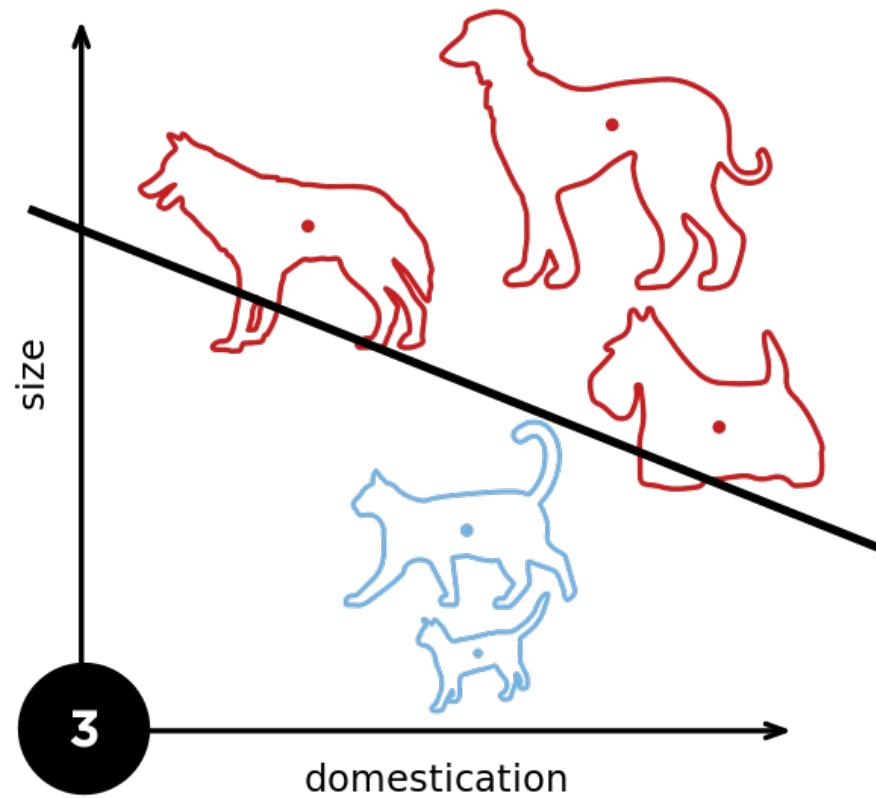


From wikipedia

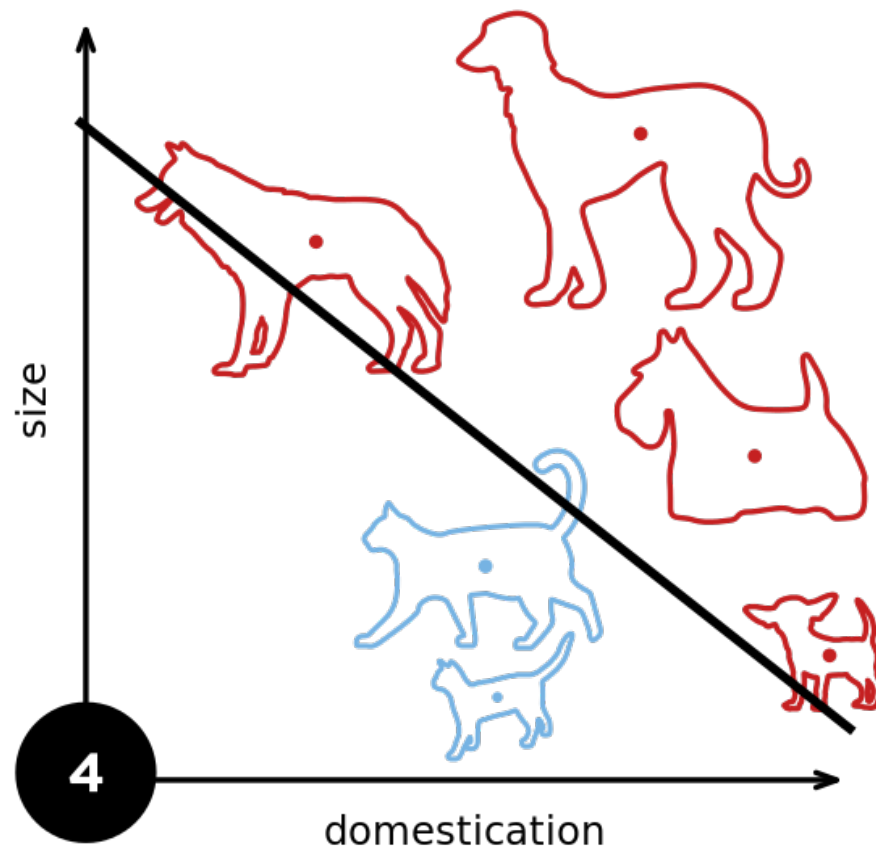




From wikipedia



From wikipedia



From wikipedia

# Convergence Theorem

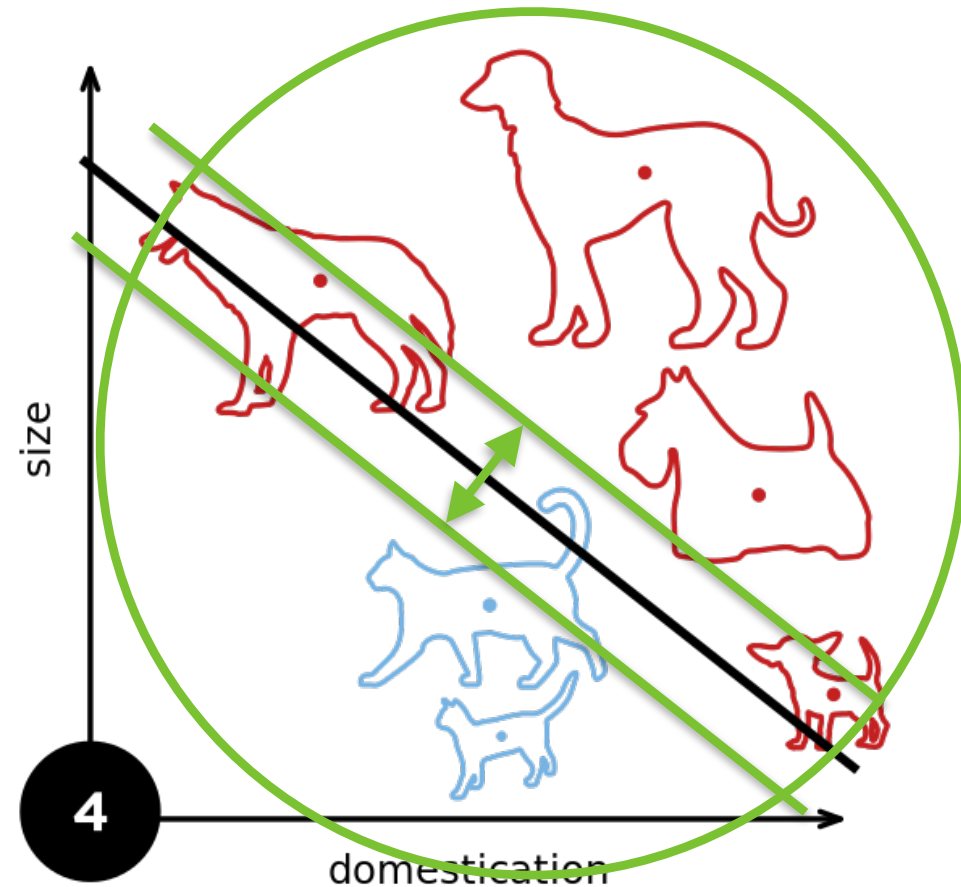
- Radius  $r$  enclosing the data
- Margin  $\rho$  separating the classes

$$y(\mathbf{x}^\top \mathbf{w} + b) \geq \rho$$

for  $\|\mathbf{w}\|^2 + b^2 \leq 1$

- Guaranteed that perceptron will converge after

$$\frac{r^2 + 1}{\rho^2} \text{ steps}$$



# Consequences

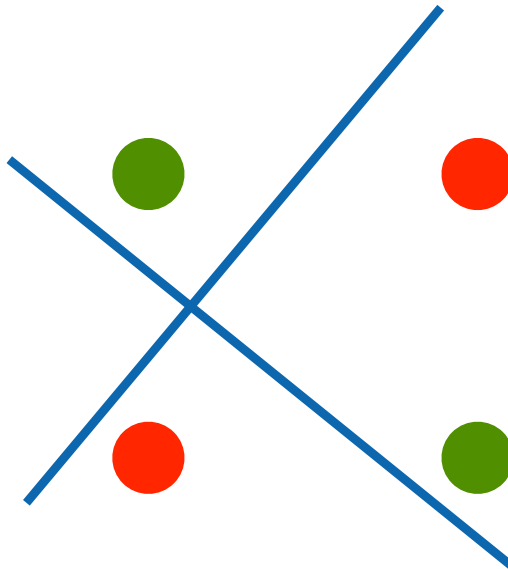
- Only need to store errors. This gives a compression bound for perceptron.
- Fails with noisy data

do NOT train your avatar with perceptrons

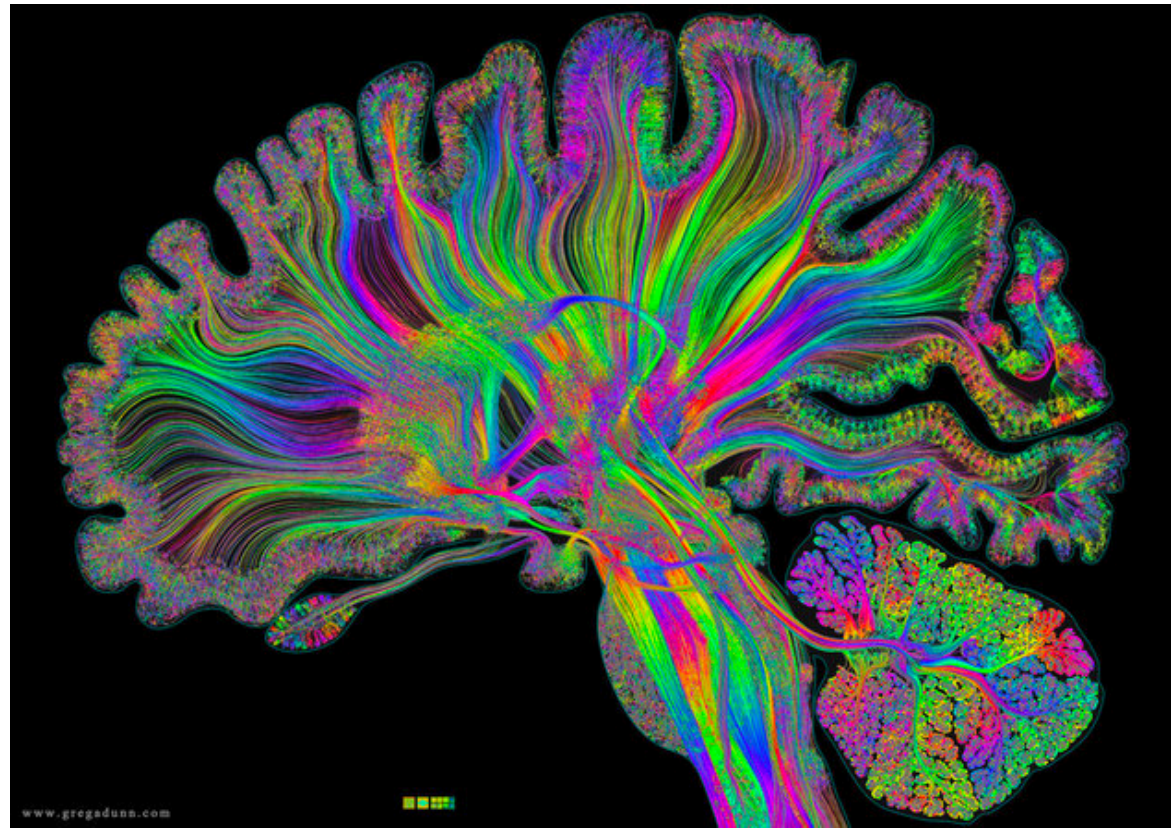


# XOR Problem (Minsky & Papert, 1969)

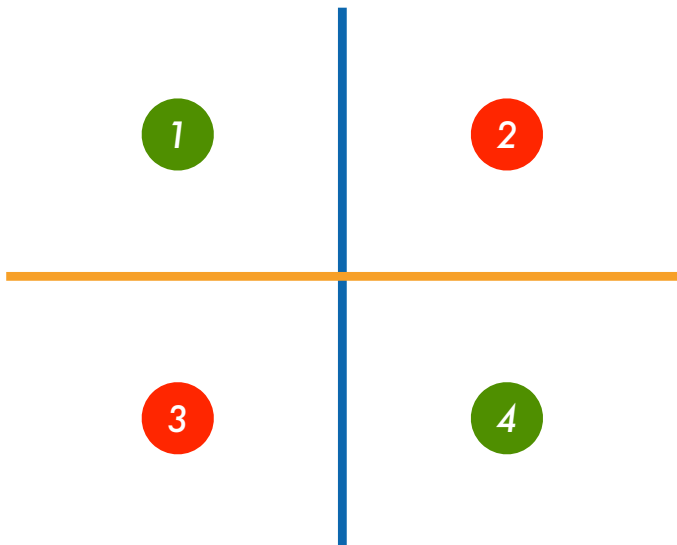
The perceptron cannot learn an XOR function  
(neurons can only generate linear separators)



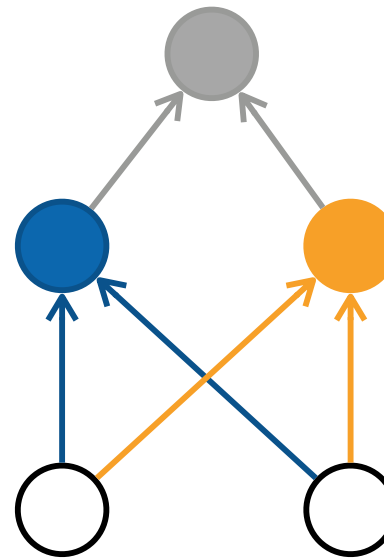
# Multilayer Perceptron



# Learning XOR

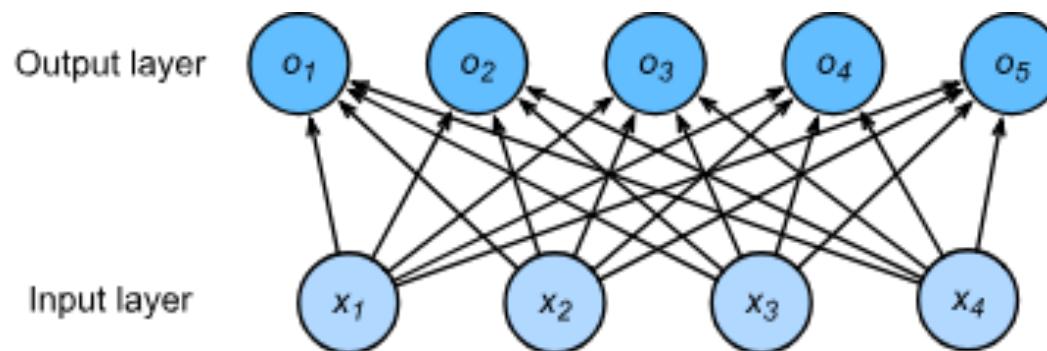


	1	2	3	4
	+	-	+	-
	+	+	-	-
product	+	-	-	+

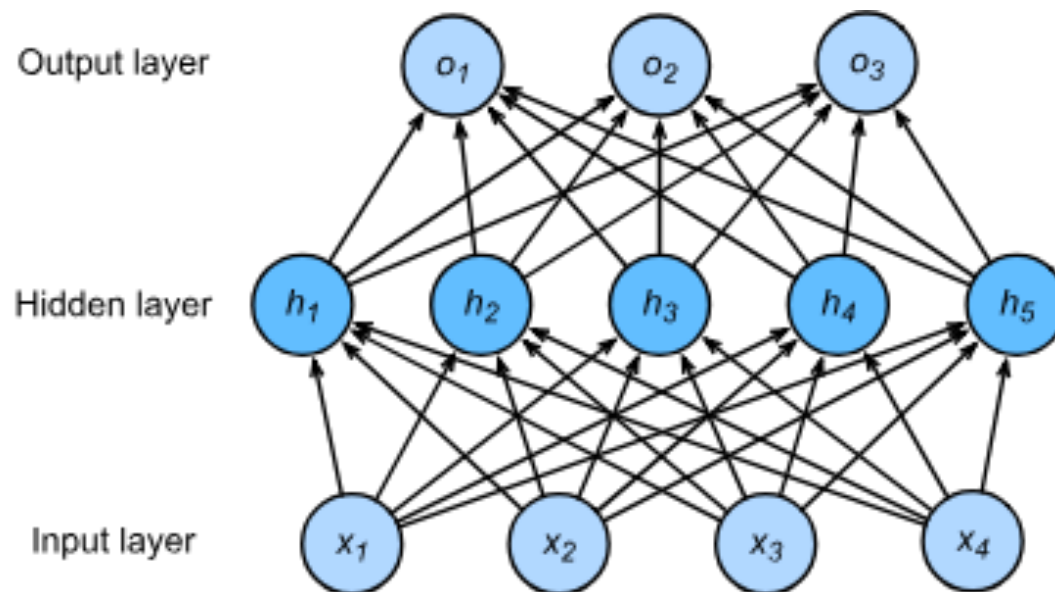




# Single Hidden Layer



# Single Hidden Layer



Hyperparameter - size  $m$  of hidden layer

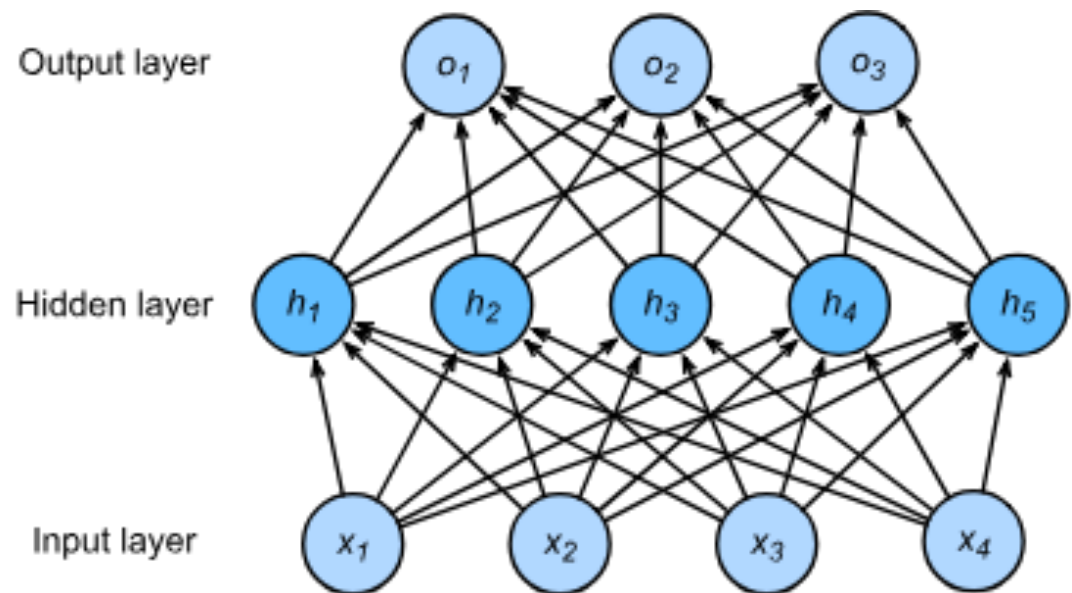
# Single Hidden Layer

- Input  $\mathbf{x} \in \mathbb{R}^n$
- Hidden  $\mathbf{W}_1 \in \mathbb{R}^{m \times n}, \mathbf{b}_1 \in \mathbb{R}^m$
- Output  $\mathbf{w}_2 \in \mathbb{R}^m, b_2 \in \mathbb{R}$

$$\mathbf{h} = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{o} = \mathbf{w}_2^T \mathbf{h} + b_2$$

$\sigma$  is an element-wise  
activation function



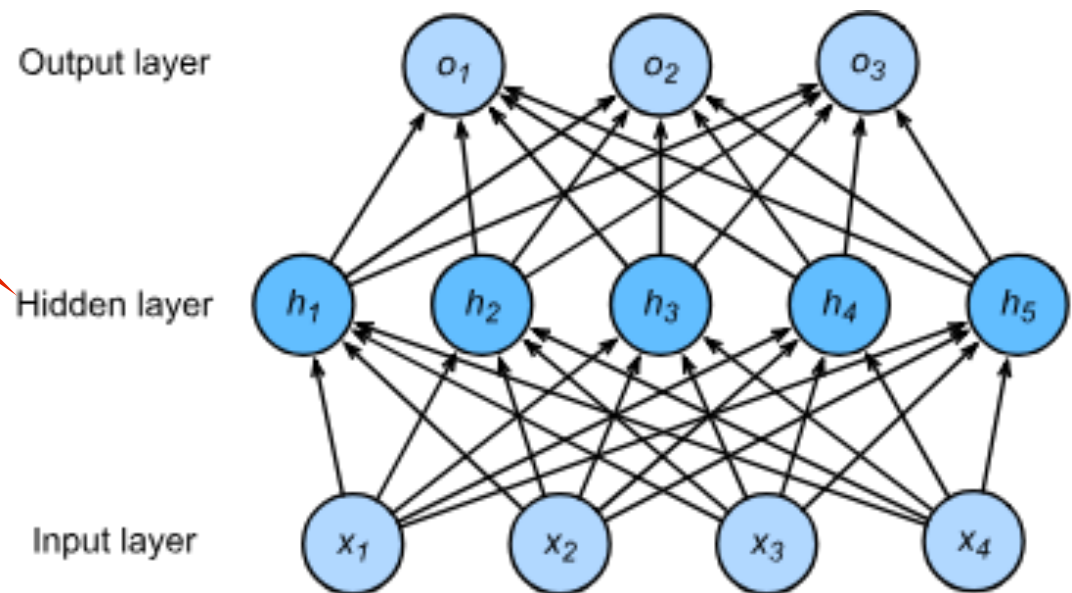
# Single Hidden Layer

Why do we need an a  
nonlinear activation?

$$\mathbf{h} = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{o} = \mathbf{w}_2^T \mathbf{h} + b_2$$

$\sigma$  is an element-wise  
activation function



# Single Hidden Layer

Why do we need an a  
nonlinear activation?

$$\mathbf{h} = \mathbf{W}_1 \mathbf{x} + \mathbf{b}_1$$

$$\mathbf{o} = \mathbf{w}_2^T \mathbf{h} + b_2$$

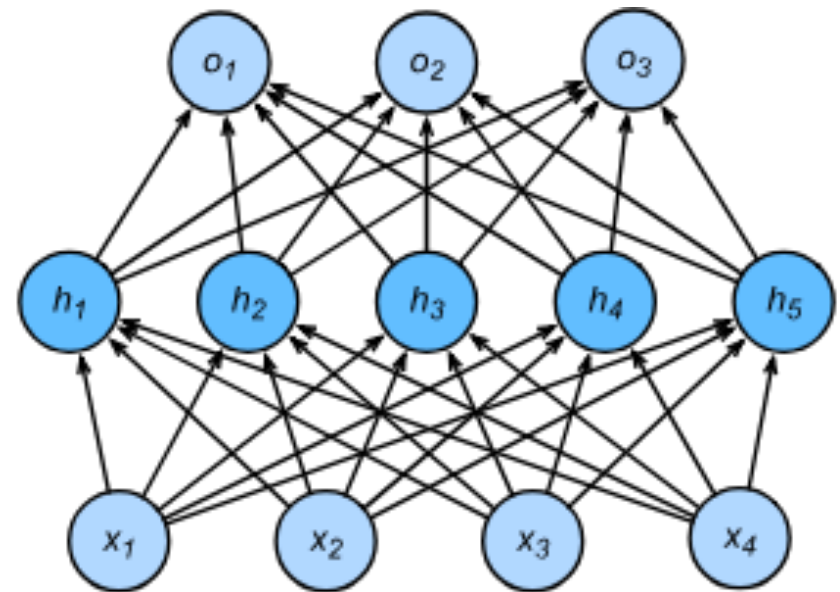
$$\text{hence } o = \mathbf{w}_2^T \mathbf{W}_1 \mathbf{x} + b'$$

Linear ...

Output layer

Hidden layer

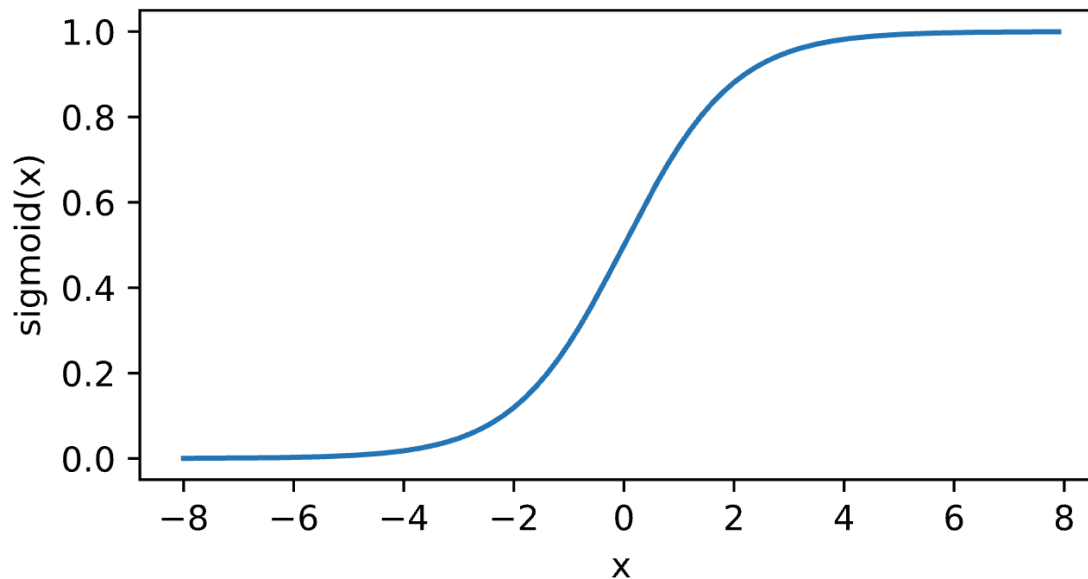
Input layer



# Sigmoid Activation

Map input into (0, 1), a soft version of  $\sigma(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$

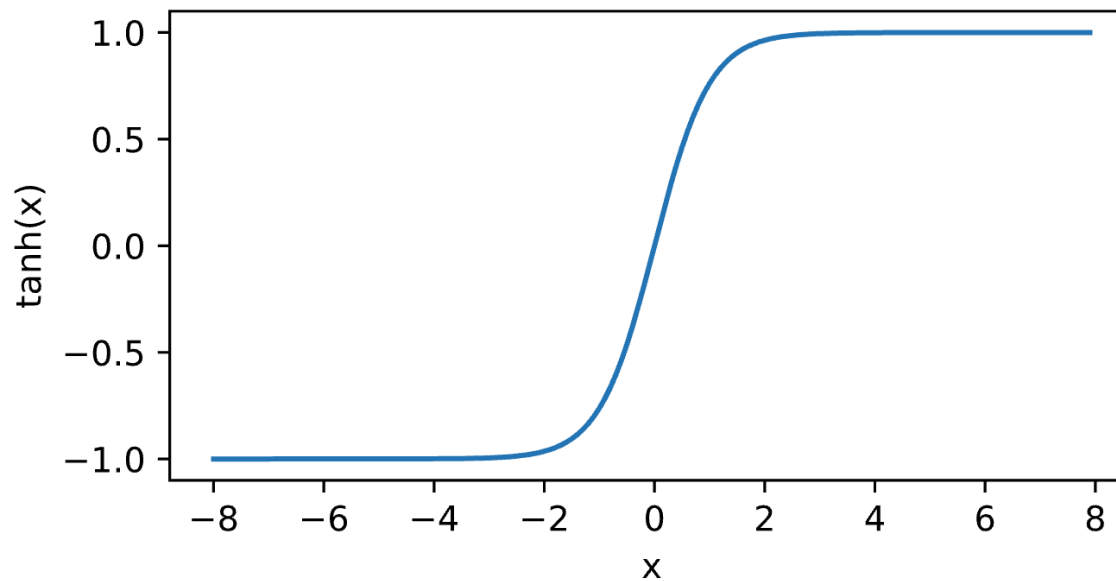
$$\text{sigmoid}(x) = \frac{1}{1 + \exp(-x)}$$



# Tanh Activation

Map inputs into (-1, 1)

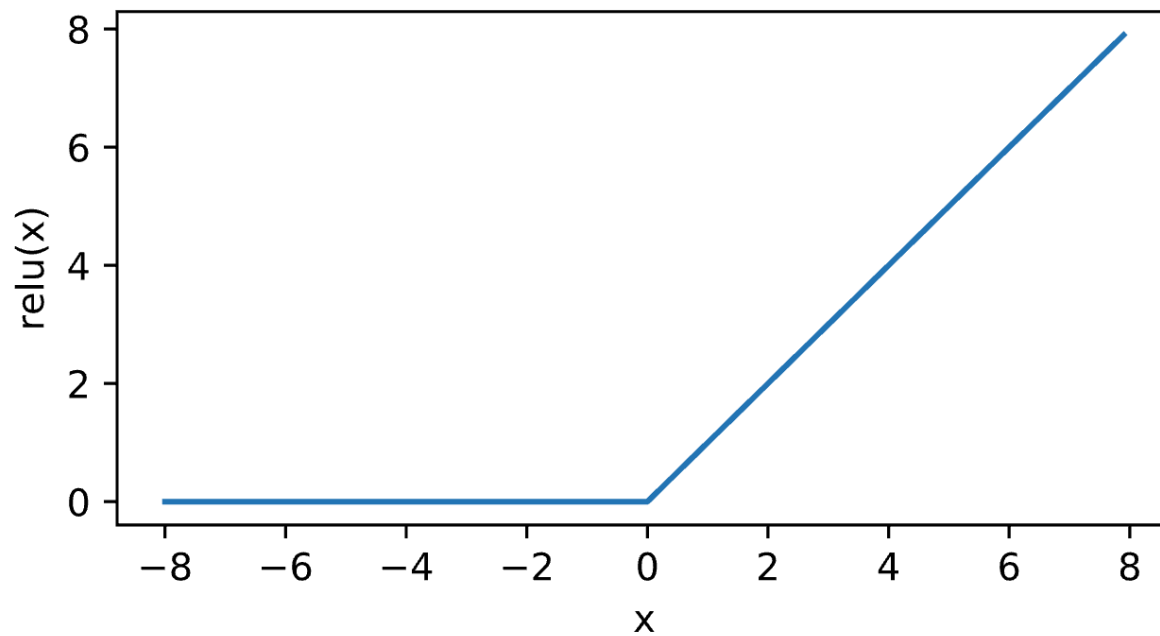
$$\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$$



# ReLU Activation

ReLU: rectified linear unit

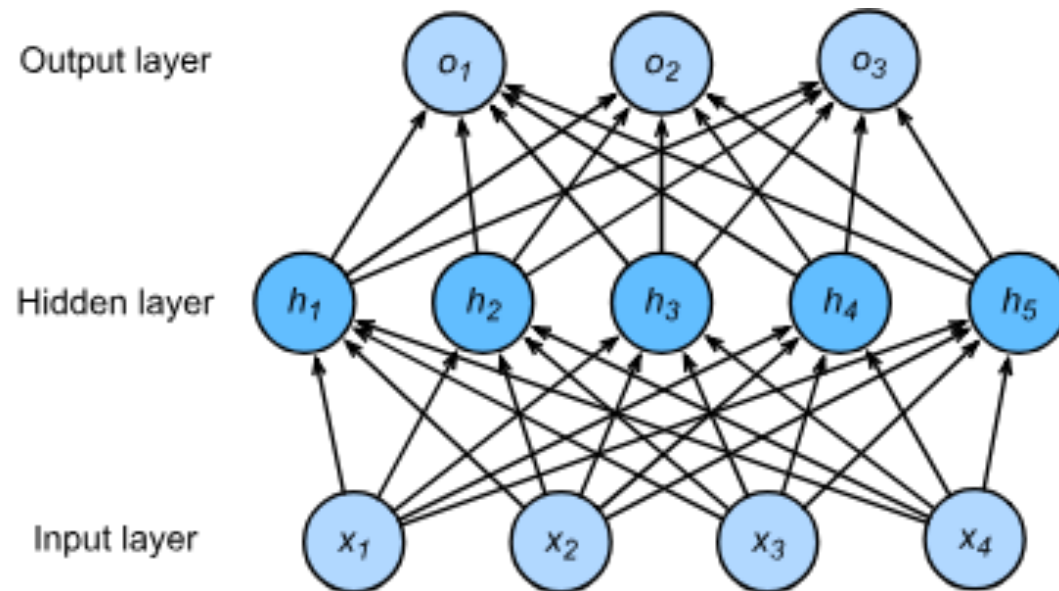
$$\text{ReLU}(x) = \max(x, 0)$$





# Multiclass Classification

$$y_1, y_2, \dots, y_k = \text{softmax}(o_1, o_2, \dots, o_k)$$



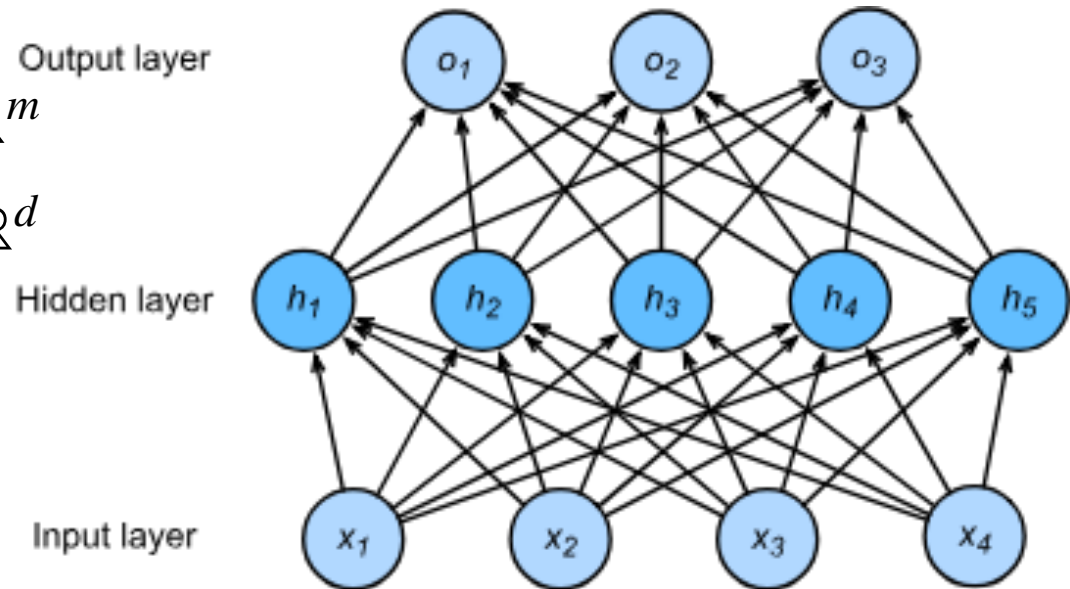
# Multiclass Classification

- Input  $\mathbf{x} \in \mathbb{R}^n$
- Hidden  $\mathbf{W}_1 \in \mathbb{R}^{m \times n}$  and  $\mathbf{b}_1 \in \mathbb{R}^m$
- Output  $\mathbf{W}_2 \in \mathbb{R}^{m \times d}$  and  $\mathbf{b}_2 \in \mathbb{R}^d$

$$\mathbf{h} = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{o} = \mathbf{w}_2^T \mathbf{h} + \mathbf{b}_2$$

$$\mathbf{y} = \text{softmax}(\mathbf{o})$$



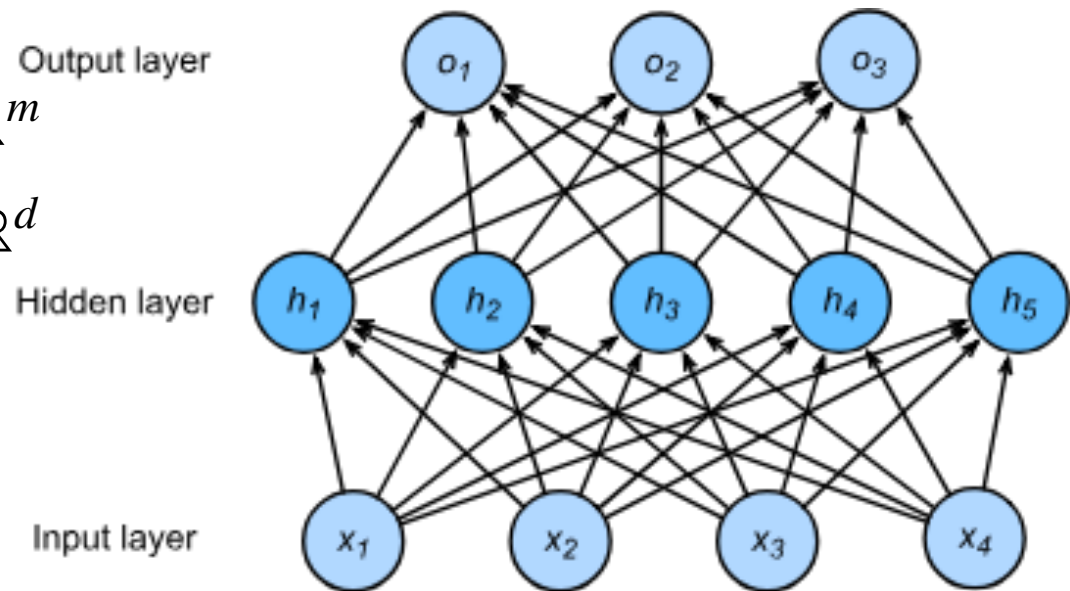
# Multiclass Classification

- Input  $\mathbf{x} \in \mathbb{R}^n$
- Hidden  $\mathbf{W}_1 \in \mathbb{R}^{m \times n}$  and  $\mathbf{b}_1 \in \mathbb{R}^m$
- Output  $\mathbf{W}_2 \in \mathbb{R}^{m \times d}$  and  $\mathbf{b}_2 \in \mathbb{R}^d$

$$\mathbf{h} = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{o} = \mathbf{w}_2^T \mathbf{h} + \mathbf{b}_2$$

$$\mathbf{y} = \text{softmax}(\mathbf{o})$$



# Multiple Hidden Layers

$$\mathbf{h}_1 = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

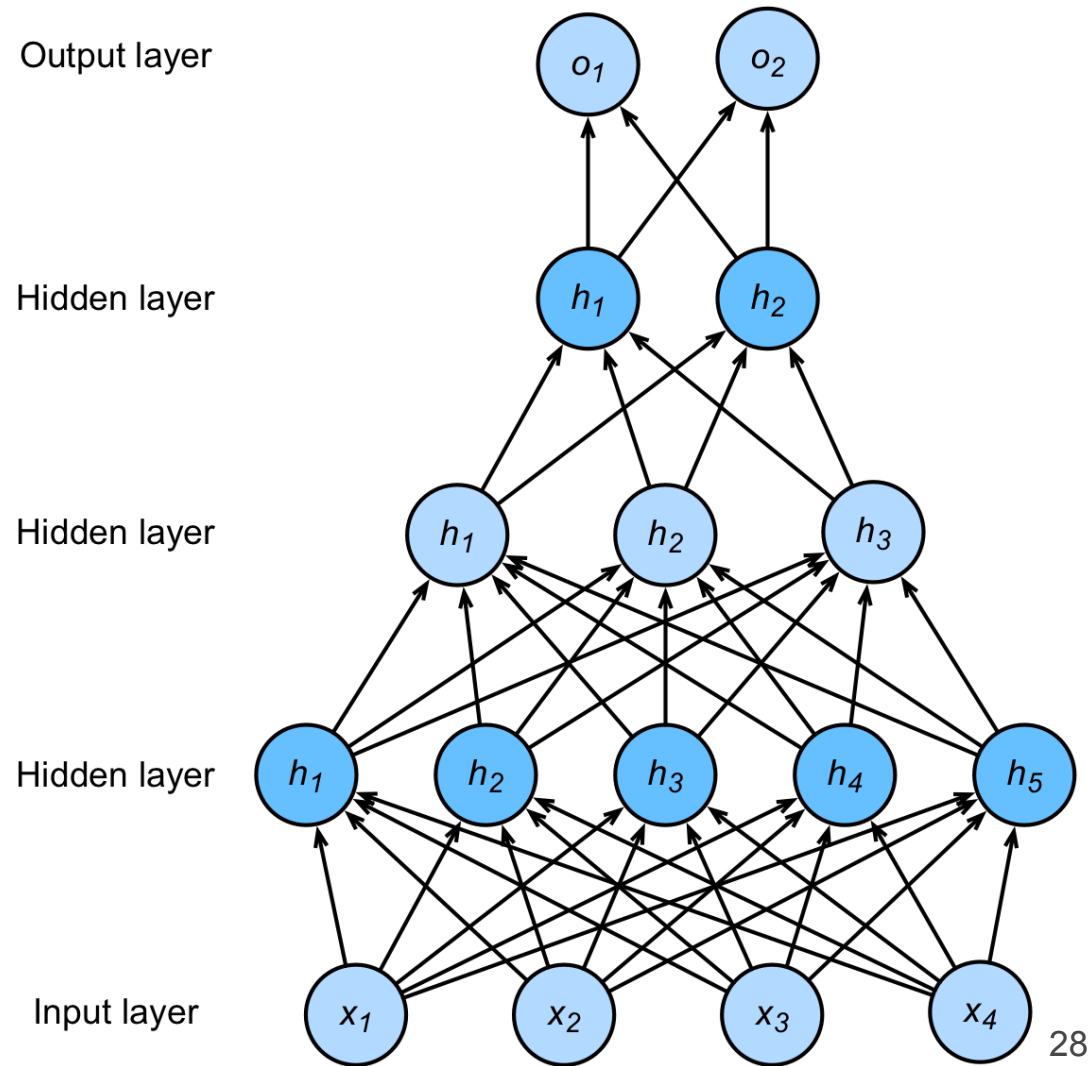
$$\mathbf{h}_2 = \sigma(\mathbf{W}_2 \mathbf{h}_1 + \mathbf{b}_2)$$

$$\mathbf{h}_3 = \sigma(\mathbf{W}_3 \mathbf{h}_2 + \mathbf{b}_3)$$

$$\mathbf{o} = \mathbf{W}_4 \mathbf{h}_3 + \mathbf{b}_4$$

Hyper-parameters

- # of hidden layers
- Hidden size for each layer



# Summary

- Perceptron
  - Simple updates
  - Limited function complexity
- Multilayer Perceptron
  - Multiple layers add more complexity
  - Nonlinearity is needed
  - Simple composition (but architecture search needed)